

Reflective Memory Kernel: Deep Dive Architecture

1. System Overview

The Reflective Memory Kernel (RMK) is a distributed system designed to provide long-term, evolvable memory for AI agents. It decouples "Computation" (LLM Inference) from "Memory" (Knowledge Graph), allowing the agent to maintain a persistent identity and knowledge base across sessions.

1.1 Core Components

- **Front-End Agent (Go/Gin):** The user-facing gateway. It manages WebSocket/REST connections, auth, and fast conversational loops. It acts as the "Consciousness," holding short-term context.
 - **Memory Kernel (Go):** The "Hippocampus." It handles ingestion, entity extraction, graph operations, and retrieval (consultation).
 - **Reflection Engine (Go):** The "Default Mode Network." It runs asynchronously to consolidate memory, find patterns, and generate insights.
 - **Knowledge Graph (DGraph):** The persistent storage layer. A distributed graph database storing all facts, entities, and relationships.
 - **Message Bus (NATS JetStream):** The nervous system. Handles high-throughput, persistent messaging between the Agent, Kernel, and Reflection Engine.
 - **AI Service (External):** An abstract inference layer (likely Python/FastAPI) that interfaces with LLMs (e.g., GPT-4, Claude 3.5 Sonnet) to generate text and extract entities.
-

2. Minute-Level Data Flow: Group Chat V2

The Group Chat feature introduces "Strict Data Isolation" into this architecture. Here is the step-by-step lifecycle of a message.

2.1 The Request (User to Agent)

1. **Input:** User sends `POST /api/chat` with body `{"message": "Hello team", "context_type": "group", "context_id": "group_123"}`.
2. **Auth:** `JWT Middleware` validates the `Bearer` token.
3. **Context Resolution (handleChat):**
 - The Agent inspects `context_type`.
 - **CRITICAL:** It determines the `namespace`.
 - If `group : namespace = "group_123"`
 - If `user : namespace = "user_<UserID>"`
4. **State Lookup:** The Agent retrieves the active `conversation` struct from memory or creates one.

2.2 The Consultation (Retrieval)

Before answering, the Agent asks the Kernel: "What do we know about this?"

1. **RPC Call:** Agent calls `MKClient.Consult()`.
2. **Payload:** `ConsultationRequest` struct is sent:

```
{  
    UserID: "user_A",  
    Namespace: "group_123", // TARGETS SHARED MEMORY  
    Query: "Hello team"  
}
```

3. **Kernel Execution:**

- o **Retrieval:** `ConsultationHandler` executes a DGraph query using the `namespace` filter: `func: type(Node) @filter(eq(namespace, "group_123") AND ...)`
- o **Isolation:** This query is physically incapable of seeing nodes with `namespace="user_A"`.

4. **Synthesis:** Kernel summarizes found facts into a `SynthesizedBrief`.

5. **Response:** Kernel returns `ConsultationResponse` with the brief.

2.3 The Processing (LLM Inference)

1. **Construction:** Agent constructs a prompt for the AI Service:

- o "You are a helpful assistant..."
- o "Context: [SynthesizedBrief from Group Memory]"
- o "User: Hello team"

2. **Inference:** `AIClient` sends this to the external AI Service.

3. **Generation:** The LLM generates a response ("Hello! How can I help the team?").

2.4 The Ingestion (Memory Formation)

Now the conversation must be stored.

1. **Event Creation:** Agent creates a `TranscriptEvent`:

```
{  
    Namespace: "group_123", // PERSISTS TO SHARED MEMORY  
    UserQuery: "Hello team",  
    AIResponse: "Hello!..."  
}
```

2. **Async Publish:** Agent publishes this event to NATS topic `transcripts.user_A` (or a group-specific topic).

3. **Kernel Ingestion:**

- o `IngestionPipeline` receives the event.
- o **Extraction:** It calls the AI Service to extract entities (`Type: Group`, `Name: Team`).
- o **Graph Write:** It writes these nodes to DGraph, explicitly setting `namespace = "group_123"` on every new Node and Edge.

3. Data Models (The Schema)

The DGraph schema is the source of truth.

3.1 The User & Group Nodes

- `User` (`type User`): Represents an identity.
- `Group` (`type Group`):
 - `uid` : Unique ID.
 - `name` : "Engineering Team".
 - `created_by` : Edge to a `User` node (Admin).
 - `group_members` : List of Edges to `User` nodes.
 - **NO MEMORY**: Validates membership only. Memory is stored in *other* nodes with the group's namespace.

3.2 The Knowledge Nodes (`type Node`)

Everything else (Facts, Entities, Events) is a `Node`.

- `namespace` (string, indexed): The security boundary. "user_123" or "group_456".
 - `activation` (float): Priority score (0.0-1.0) for retrieval. Decays over time (Reflection Engine).
 - `edges` : Relationships (`likes` , `knows` , `works_on`). All edges are directed.
-

4. Frontend Architecture (The App Shell)

The frontend is a lightweight Single Page Application (SPA).

- **Global State:** `currentContext` (`{type: 'group', id: '...'}`).
- **Navigation:**
 - Clicking "My Memory" sets `currentContext = {type: 'private'}` .
 - Clicking "Group X" sets `currentContext = {type: 'group', id: 'X'}` .
- **API Layer:** `sendMessage` reads `currentContext` and injects it into the JSON payload of every request.

5. Security Model (Subuser & Admin)

- **Authentication:** JWT (Stateless). User identity is confirmed per request.
 - **Admin Check:**
 - `DELETE /groups/{id}` checks if `RequesterID == Group.created_by`.
 - **Subuser Creation:**
 - Admin sends `POST /subusers`.
 - System creates a new `User` node (detached from any real email/auth provider for now).
 - System links `Group -> group_members -> NewUser`.
 - NewUser allows an AI agent (or team member) to access the `group_123` namespace.
-

6. Technical Implementation Specifications

6.1 API Endpoints (Agent Layer)

Chat Interaction

- `POST /api/chat`
 - **Payload:** `{"message": string, "context_type": "user"|"group", "context_id": string}`
 - **Handler:** `handleChat` (in `server.go`)
 - **Logic:** Resolves `namespace` -> Calls `MKClient.Consult` -> Calls `AIClient.Generate` -> Streams to NATS.

Group Management

- `GET /api/groups`
 - **Returns:** List of groups where `RequesterID` is in `group_members`.
- `POST /api/groups`
 - **Payload:** `{"name": string}`
 - **Logic:** Creates `Group` node, links `created_by` to Requester, adds Requester to `group_members`.
- `DELETE /api/groups/{id}`
 - **Auth:** Requires `RequesterID == Group.created_by`.
 - **Effect:** Deletes Group node. *Note: Does not currently cascade delete memory nodes for safety.*

Member Management

- POST /api/groups/{id}/subusers
 - Payload: { "username": string, "password": string }
 - Logic: Registers new User -> Adds to Group.
- DELETE /api/groups/{id}/members/{username}
 - Logic: Removes edge Group -> group_members -> TargetUser .

6.2 Key Code Modifications

internal/agent/server.go

- Updated ChatRequest struct: Added ContextType and ContextID .
- Updated handleChat : Logic to switch namespace variable based on context.

internal/kernel/ingestion.go

- Updated IngestionPipeline.Process : Now accepts namespace argument.
- Updated processBatchedEntities : Uses namespace to construct the DGraph namespace predicate for every new node.

```
node.Namespace = event.Namespace // e.g. "group_123"
```

internal/kernel/consultation.go

- Updated ConsultationHandler.Handle : Extracts Namespace from request.
- Updated getUserKnowledge :
 - Old: func: has(name) (implied user scope)
 - New: func: eq(namespace, \$namespace) (Explicit scope)

6.3 DGraph Schema Changes (internal/graph/schema.go)

- Added: namespace: string @index(exact) .
- Updated: TranscriptEvent and ConsultationRequest structs in Go now map to this schema field.